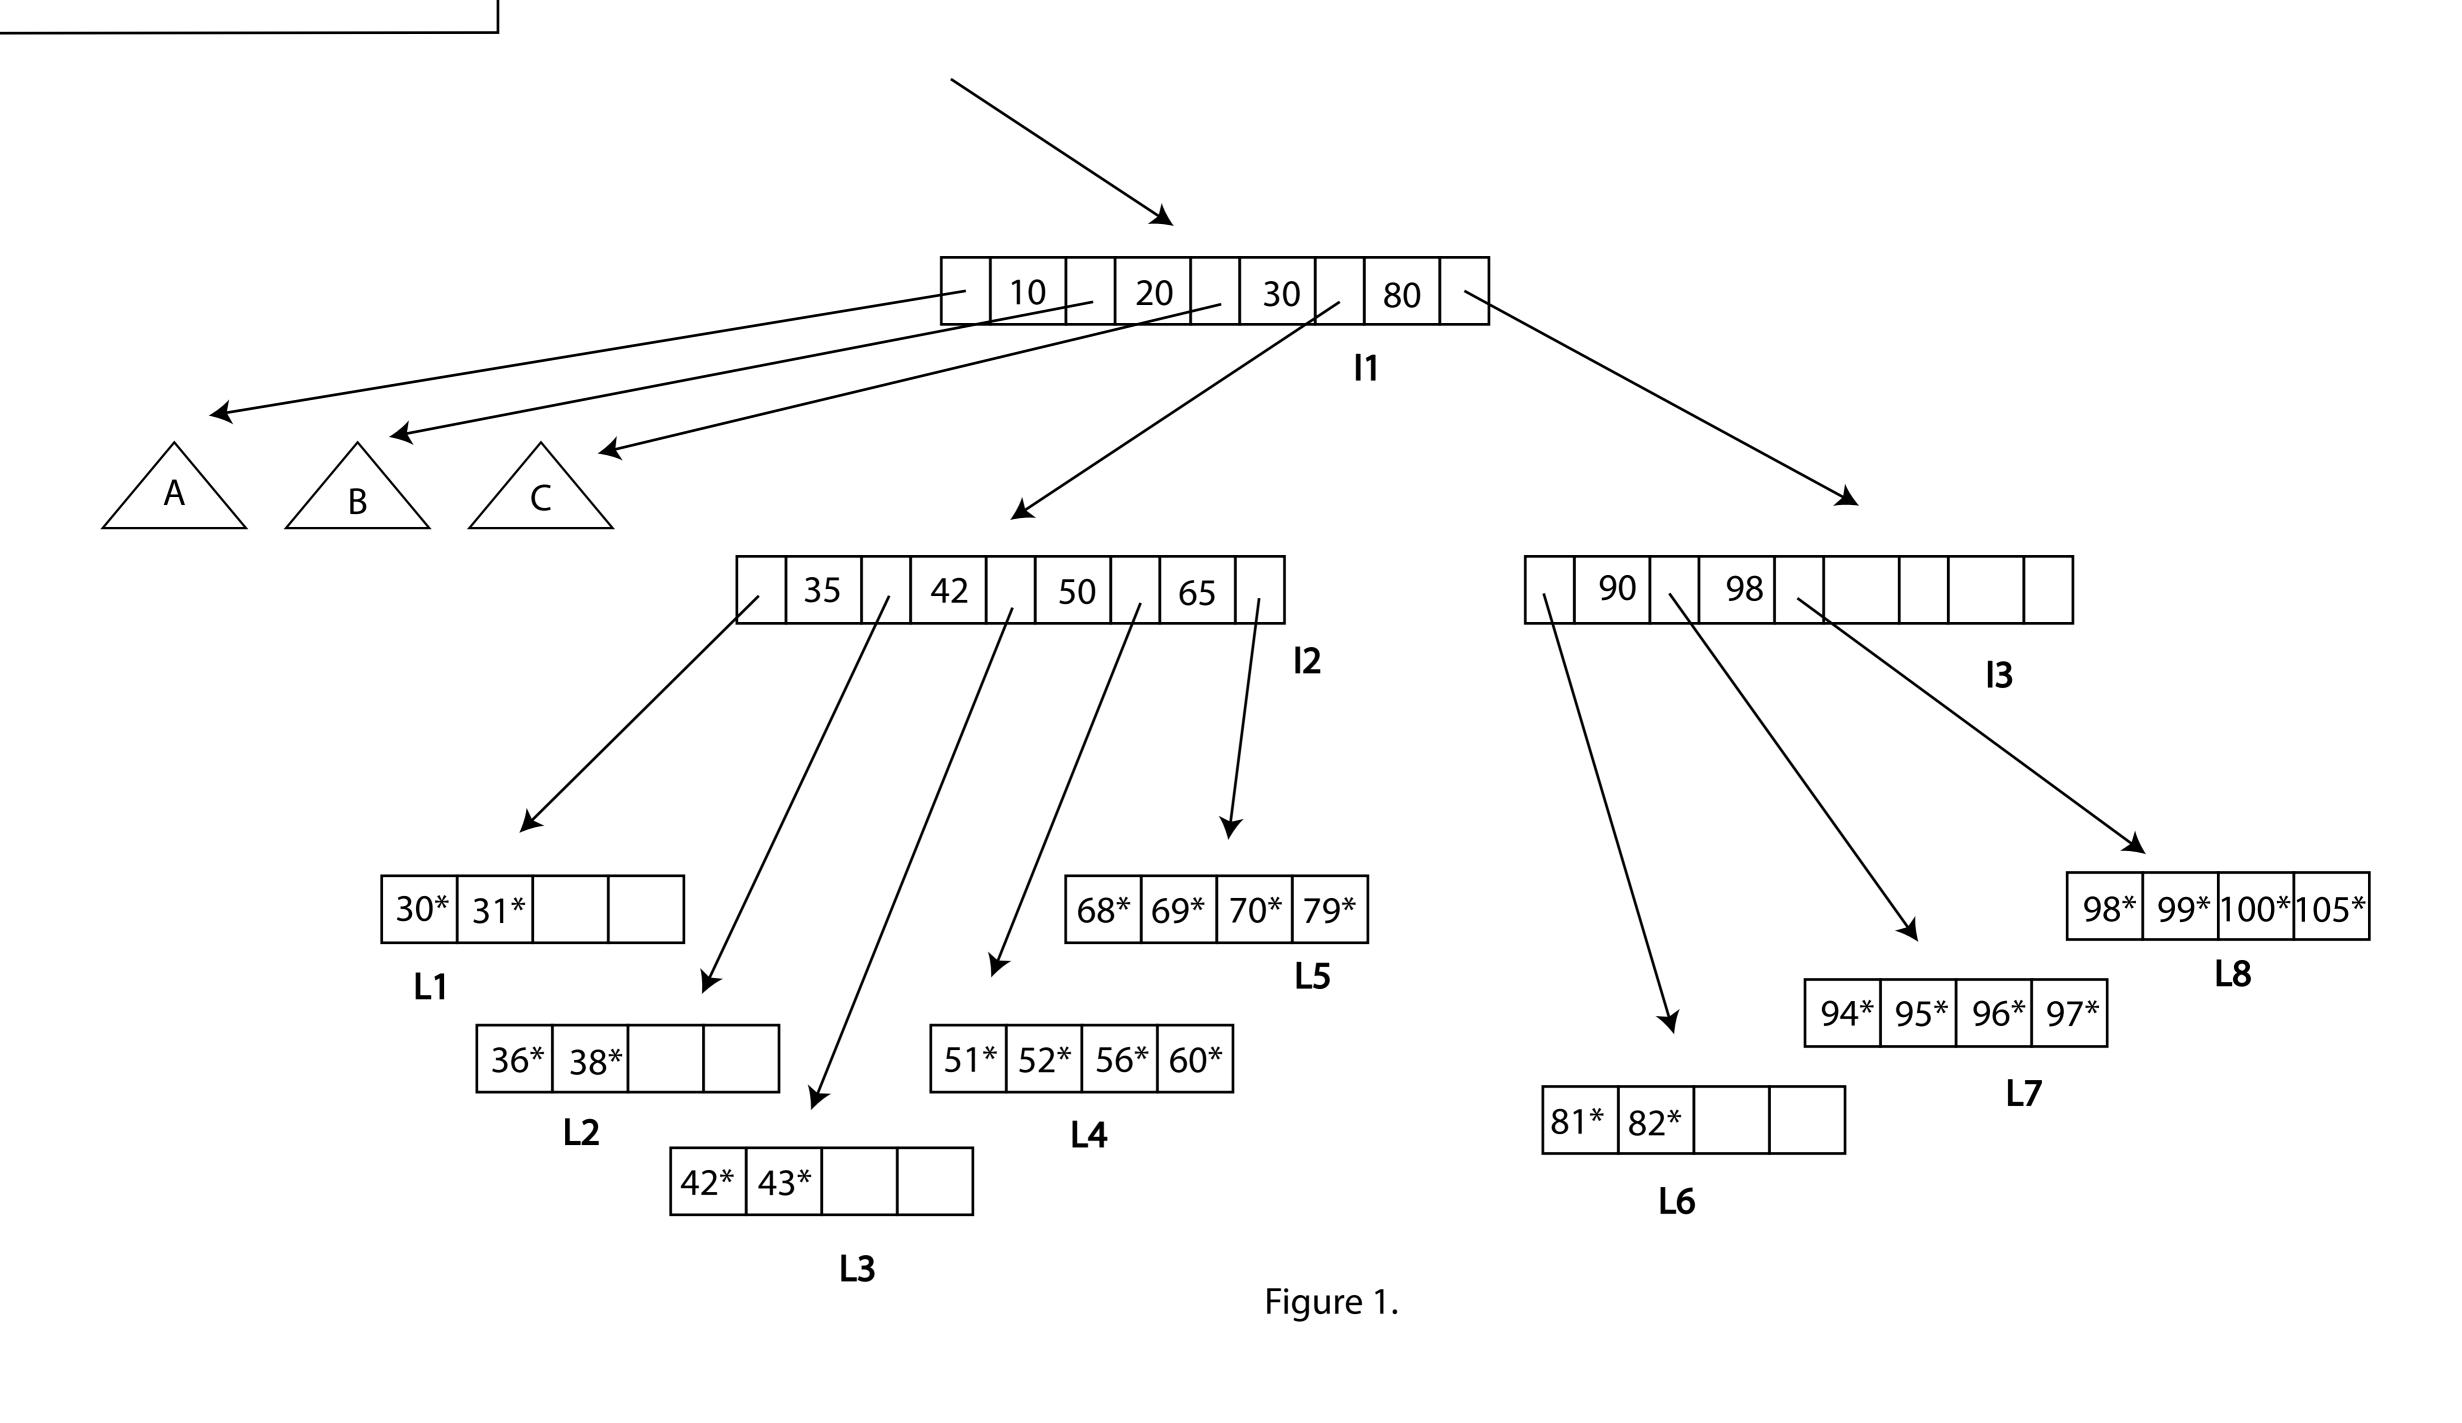
Question 1 (8 points): Consider the B+ tree index shown in Figure 1. Each intermediate node can hold up to five pointers and four key values. Each leaf can hold up to four records, and leaf nodes are doubly linked as usual, although these links are not shown in the figure. Answer the following questions. Show every step of the insertion and explain the process.

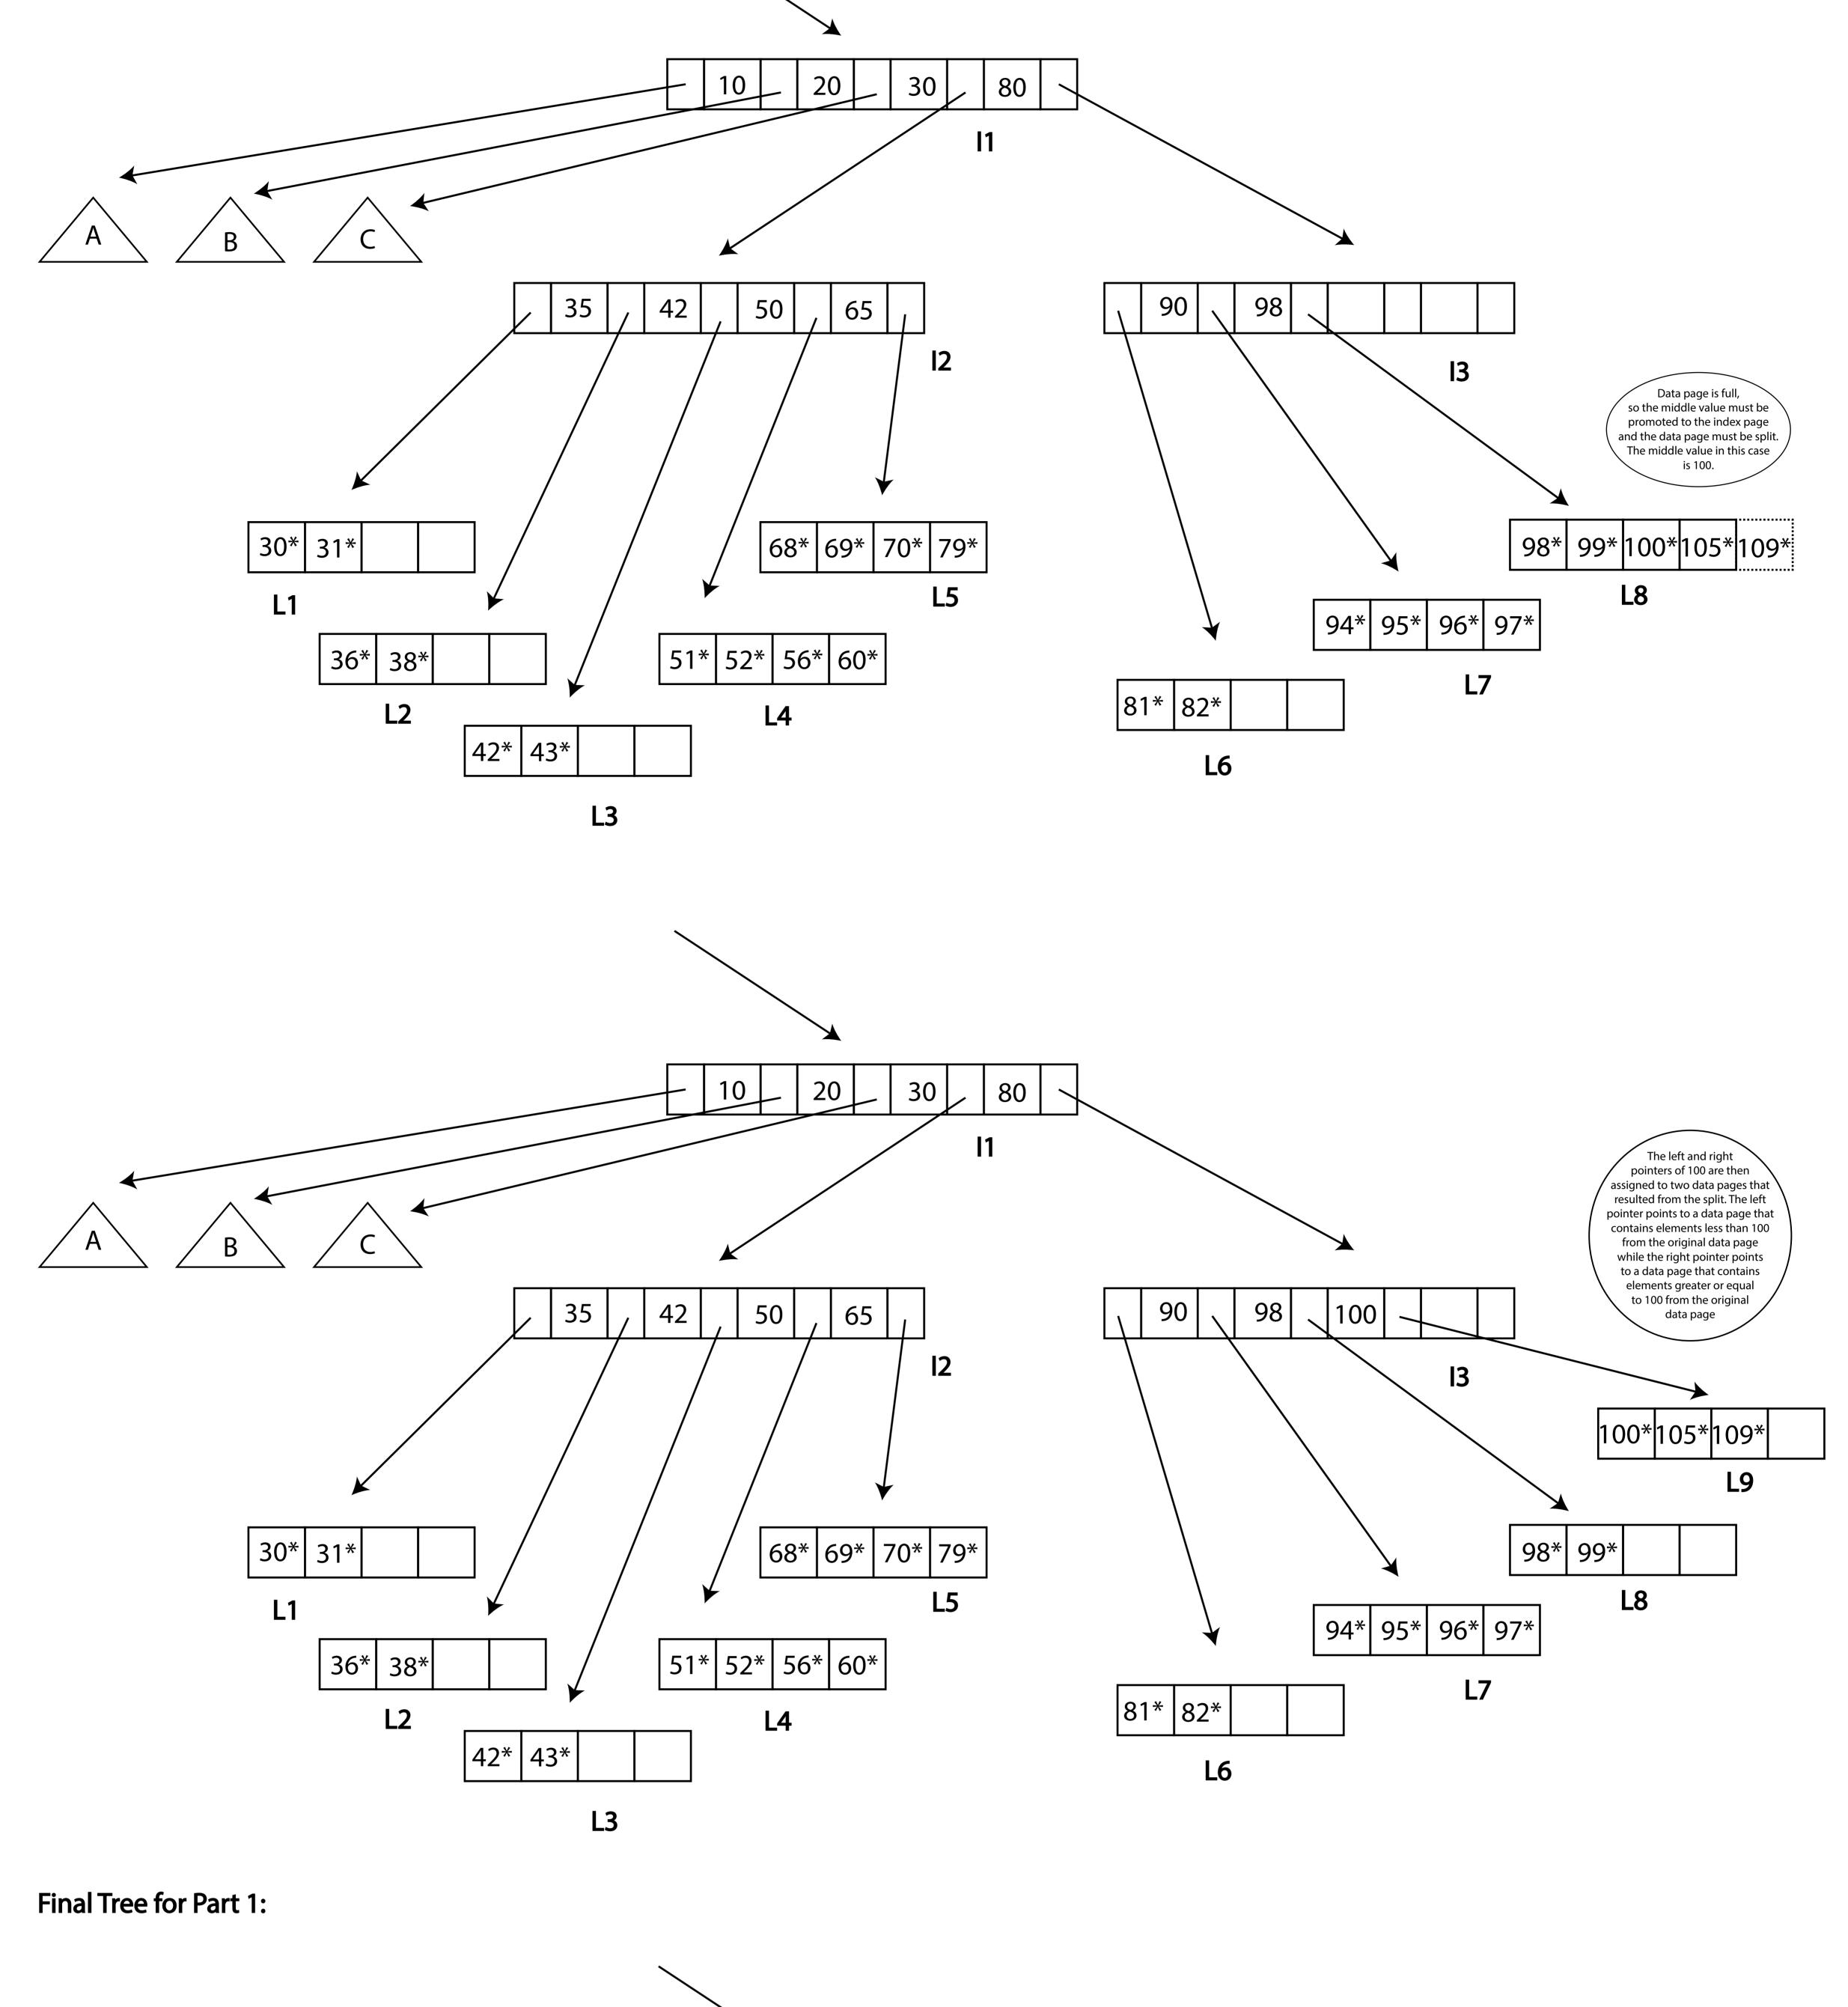


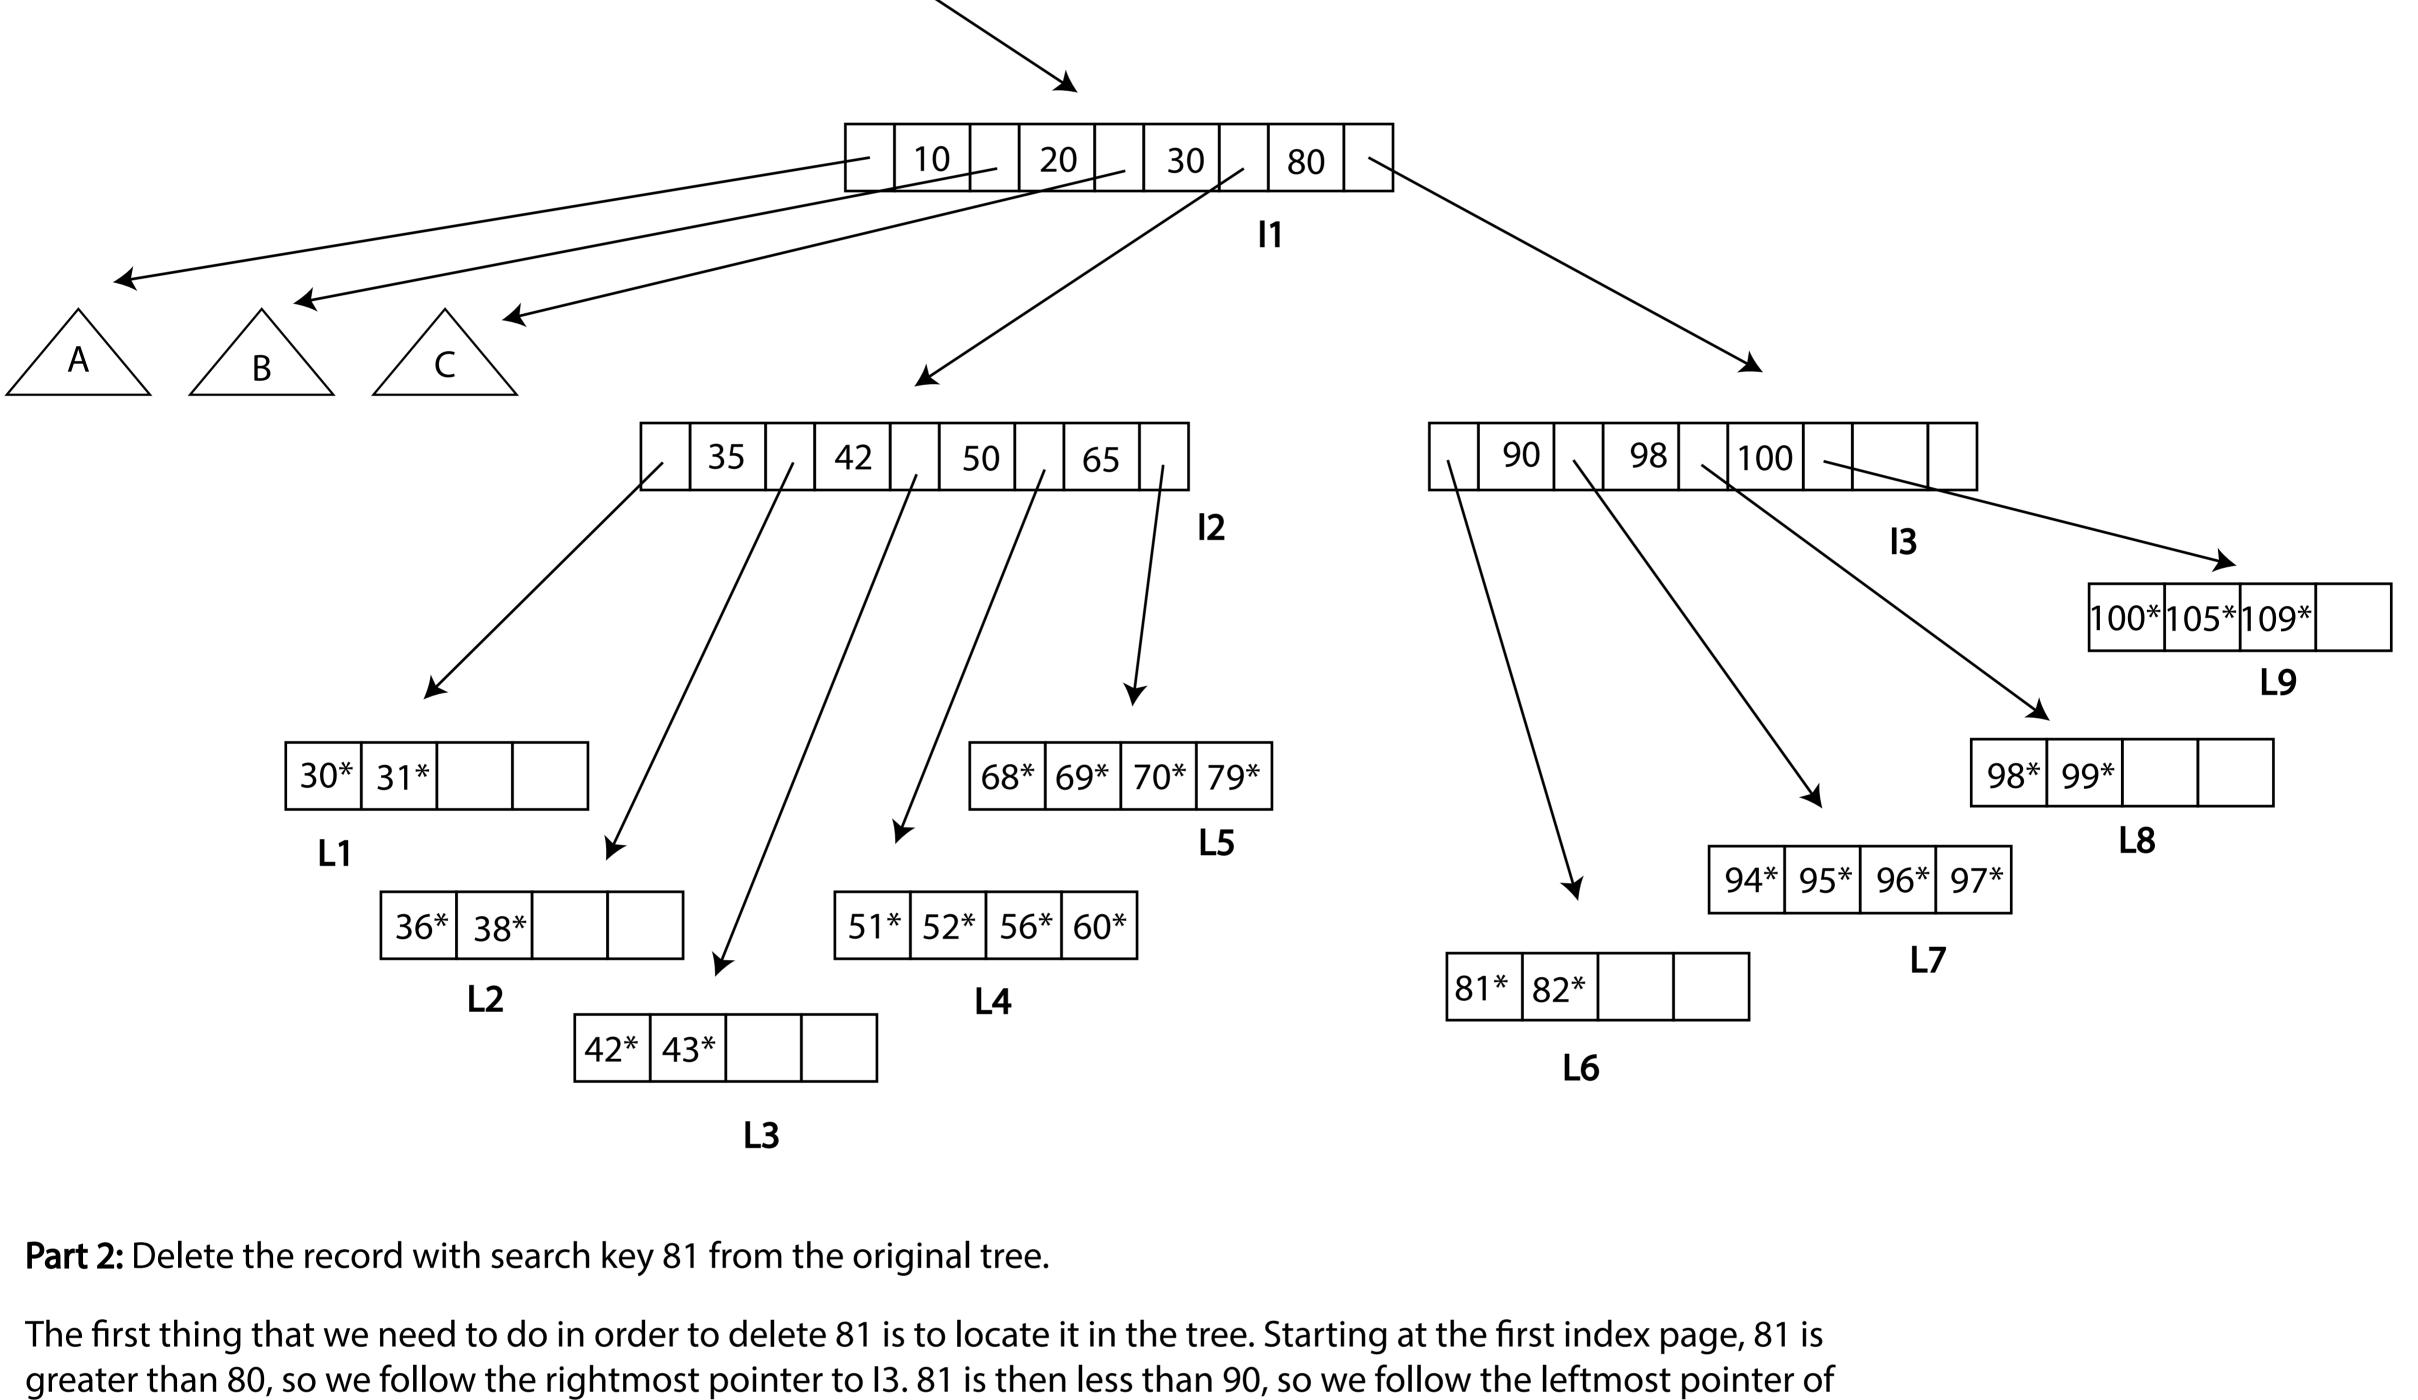
larger than 80, we follow the right pointer to the rightmost index page labeled I3. We repeat this process and find that

Part 1: Insert a record with search key 109 into the tree.

109 is larger than 98, leading us to follow the pointer to the data page L8. At this point 109 should be inserted into the data page, however the data page is full. Therefore, we need to split the data page by promoting the middle value to the index page I3 and making the left pointer of the middle value lead to a data page containing the elements smaller than the middle value and its right pointer lead to a data page containing the elements greater than or equal to the middle value. This can be done because the index page I3 is not full and promotion is thus possible.

The first thing that we need to do in order to insert 109 is to compare it to the values in the first index page. Since 109 is





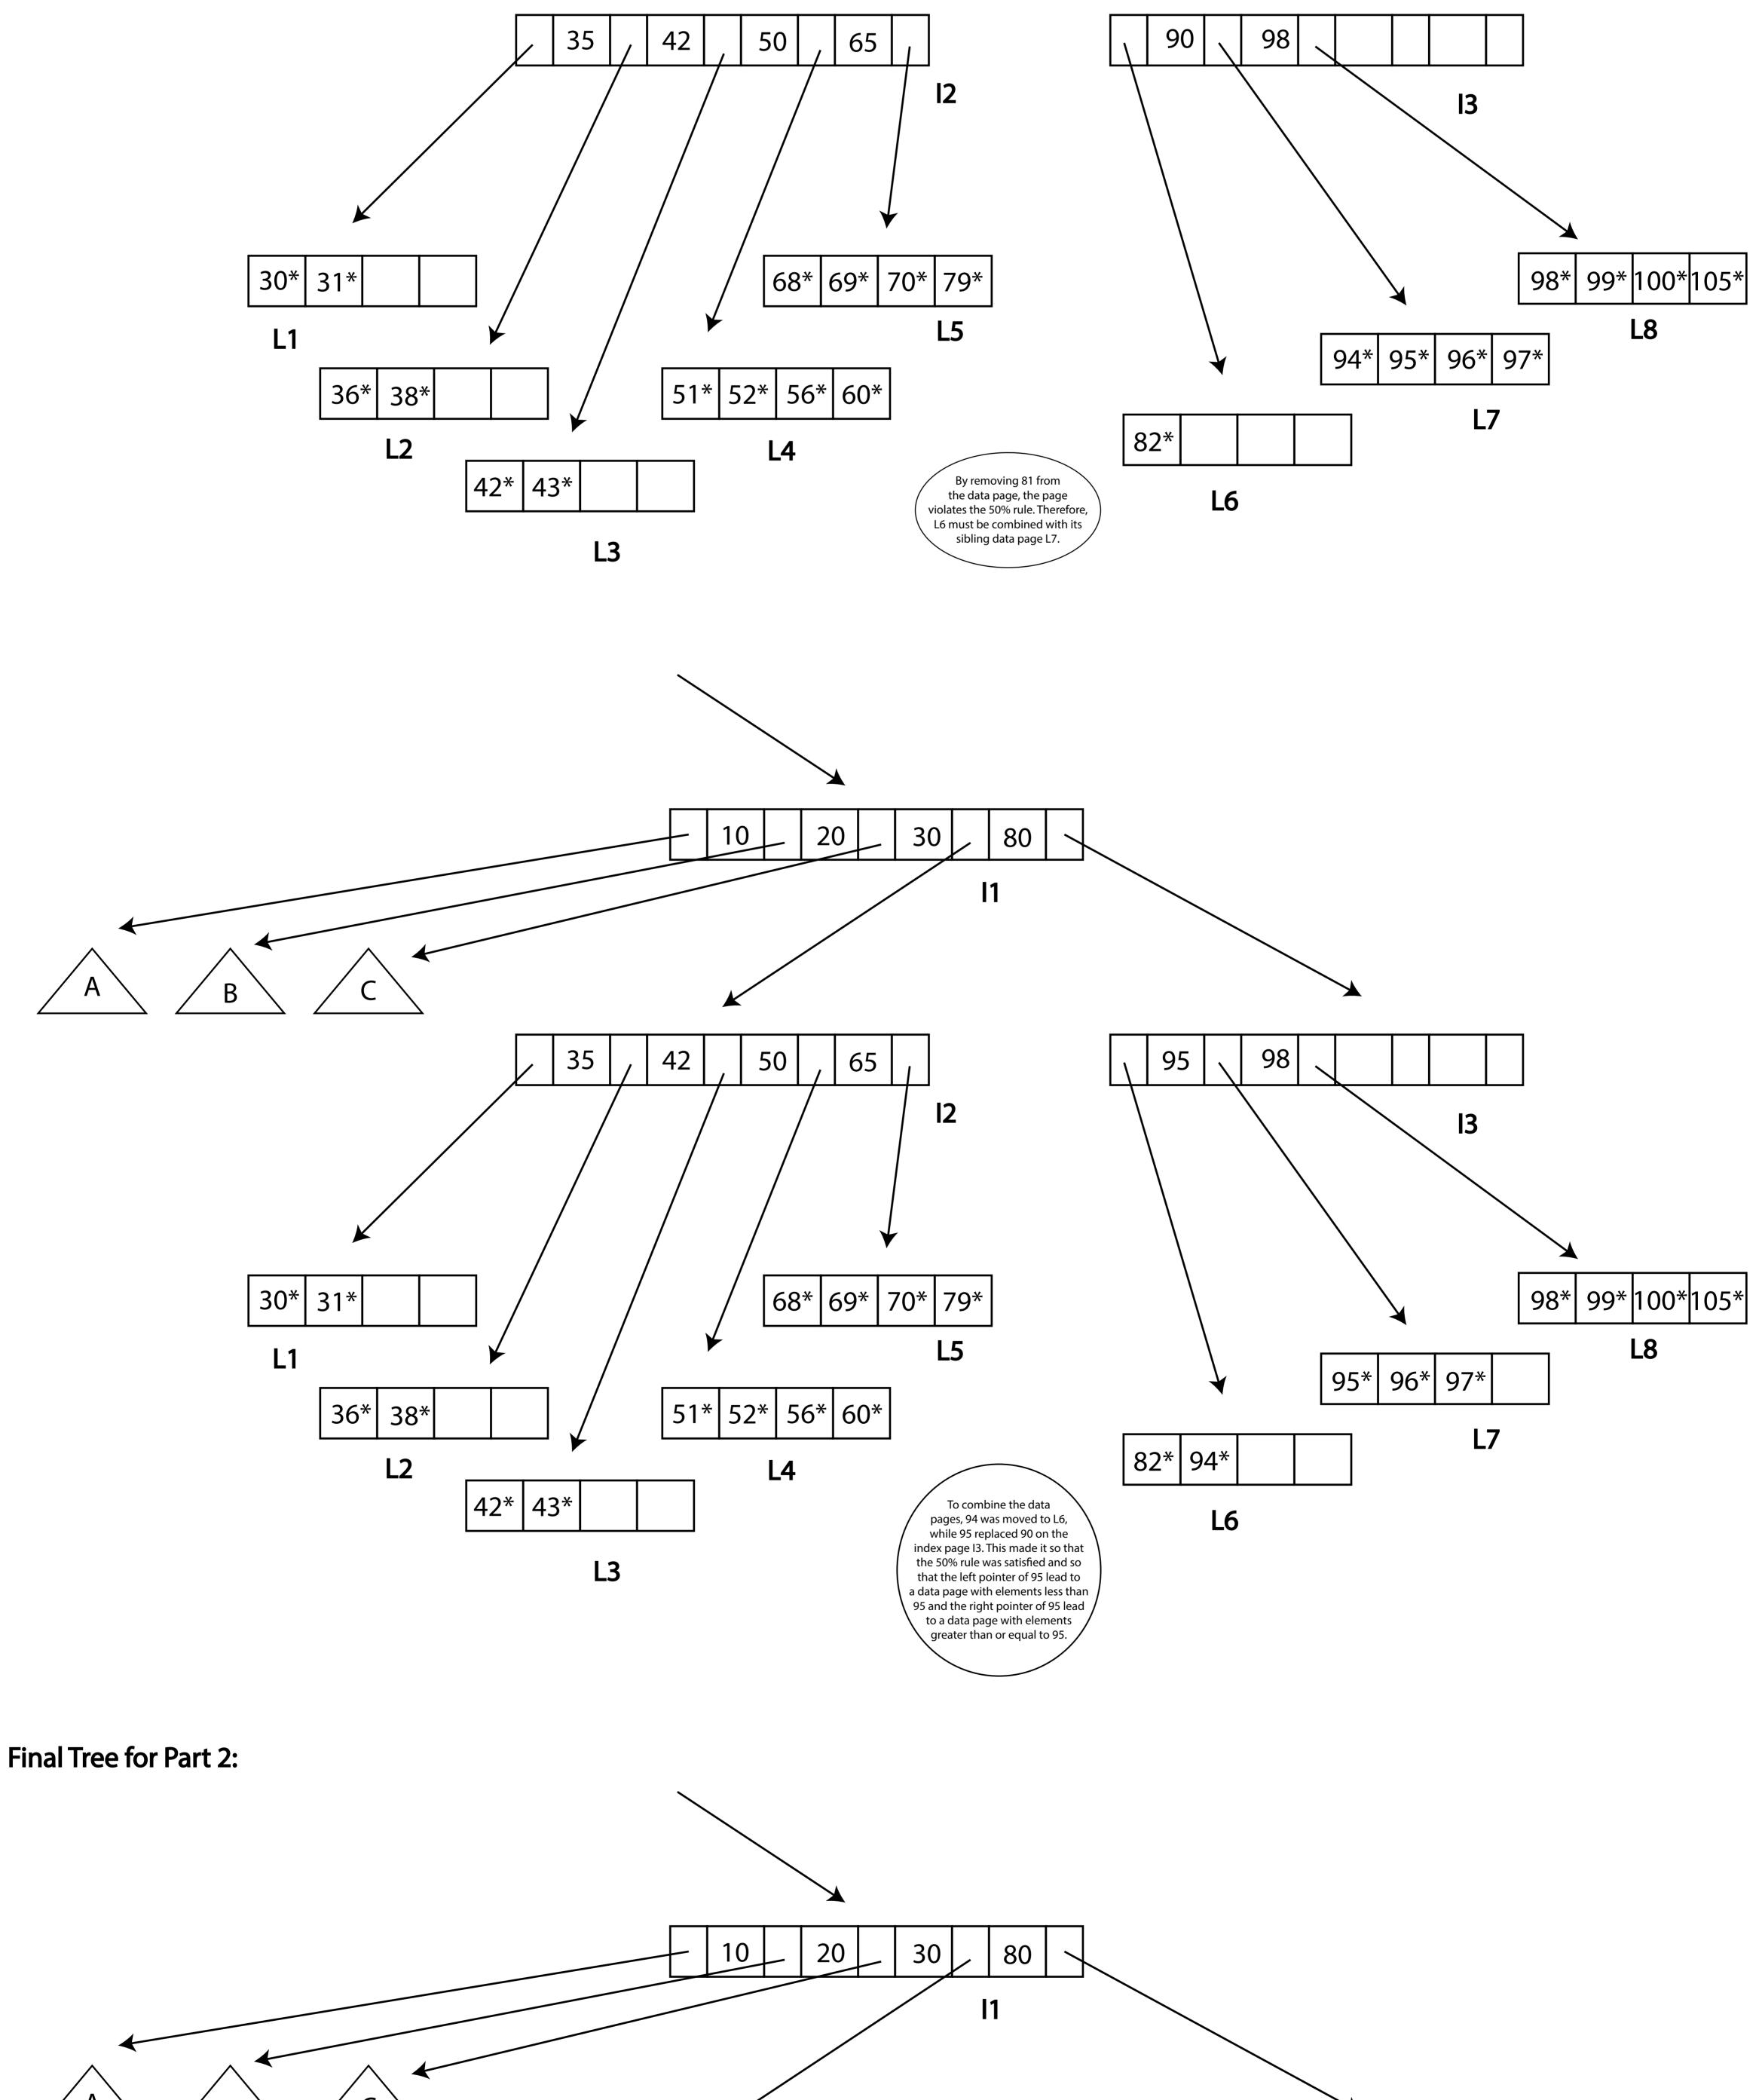
95 on index page I3.

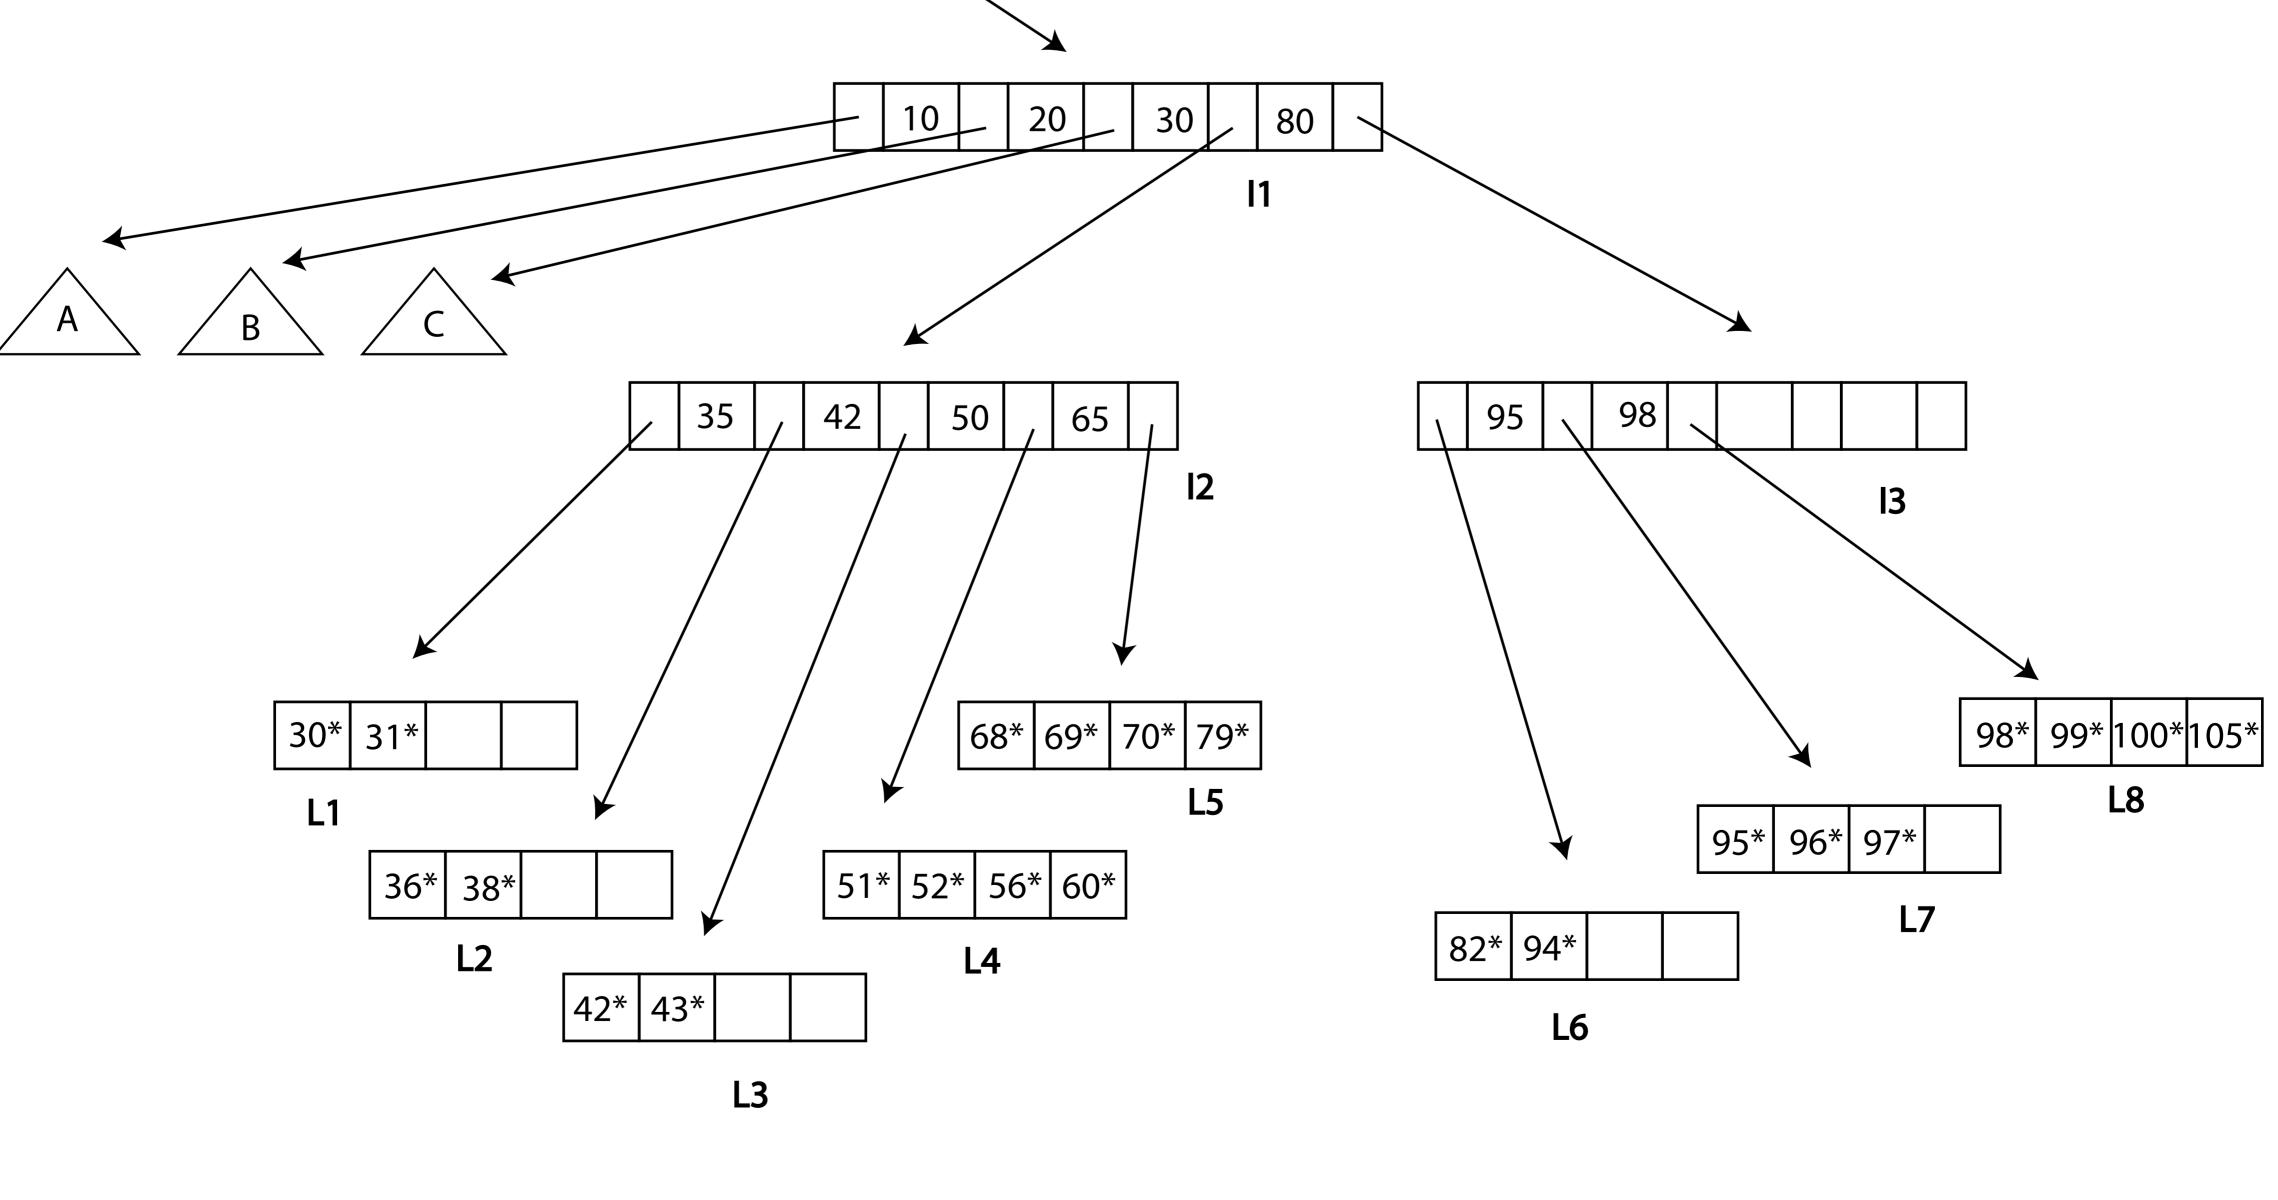
30 80 11

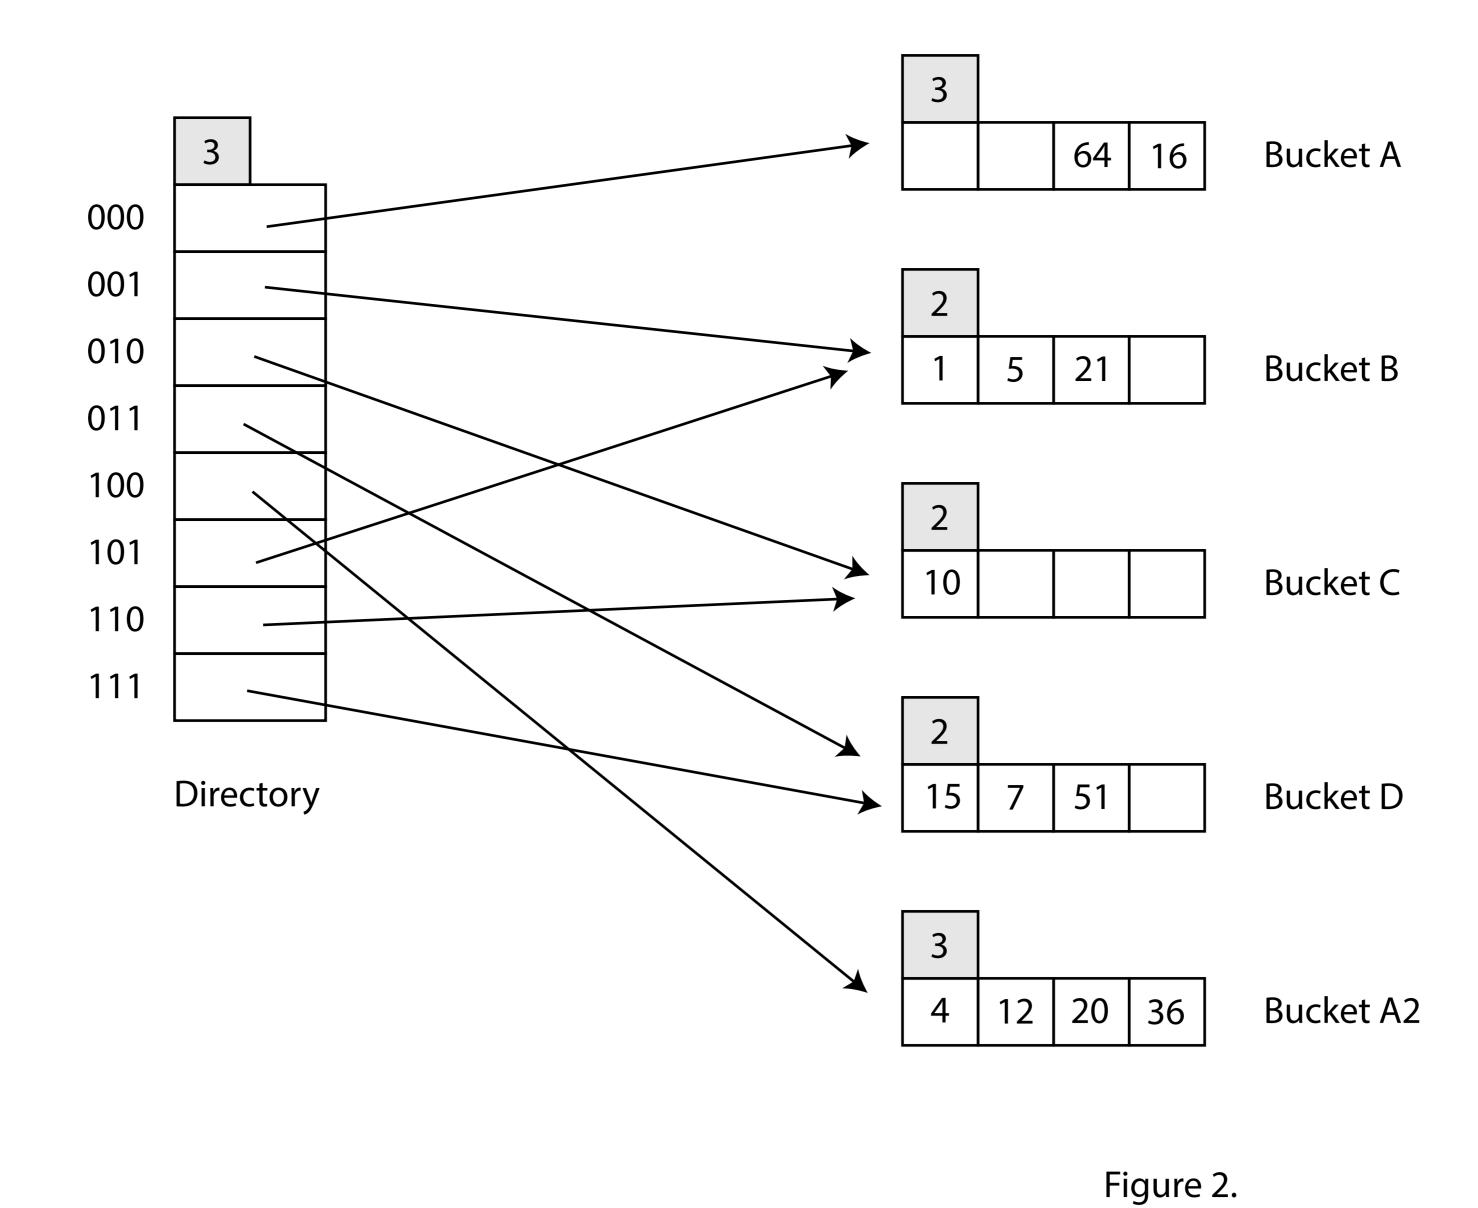
I3 to the data page L6. The data page contains two elements, 81 and 82. By removing 81 from the data page, the page

would violate the 50% rule. This means that the data page must be merged with its sibling and the index page must

be changed to reflect that merge. In this case, L6 must be merged with L7 by moving 94 to L6 and replacing 90 with







The first thing that we need to do in order to insert 68 is to find its binary representation. The binary representation of

64

Bucket A

depth, the global depth also needs to be increased and the directory size must be doubled. Bucket A2 must now be split into two buckets A2 and A3 with local depths of 4 and the contents of the old bucket A2 must be split among these two new buckets. The directory pointers must then be filled in by pointing to the correct buckets.

68 = 1000100

Final Index:

68 is 1000100, which maps to bucket A2 because the global depth is 3 and the right-most three digits are 100. Bucket

A2 is currently full, so we need to increase the local depth of the bucket to 4. Since this depth is larger than the global

## Final Index:

